

Improving Data Reuse in NPU On-chip Memory with Interleaved Gradient Order for DNN Training

Jungwoo Kim, Seonjin Na, Sanghyeon Lee, Sunho Lee, and Jaehyuk Huh

56th IEEE/ACM International Symposium on Microarchitecture (MICRO)

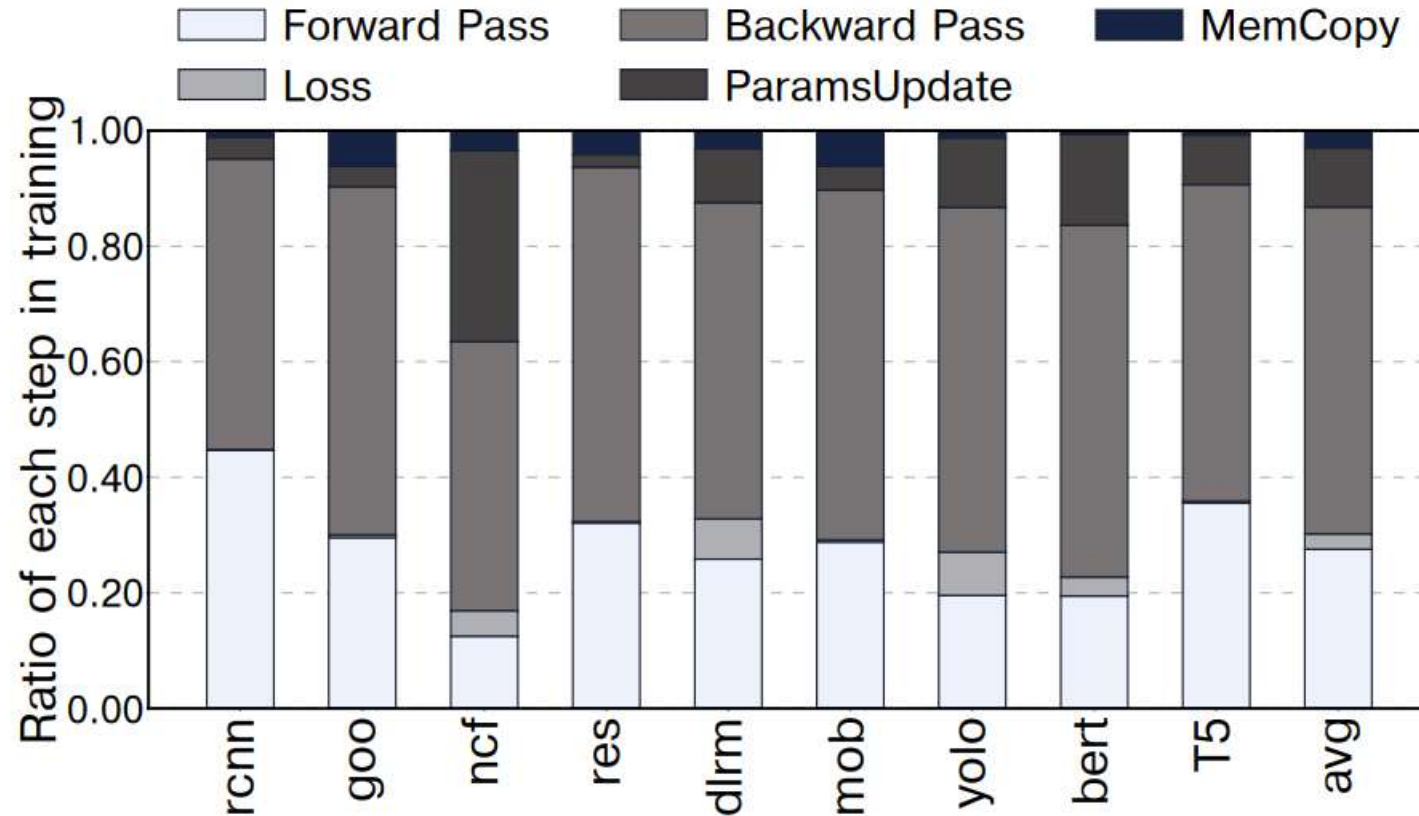
DNN Training Problem : Backward Pass!

- Computation of gradients in the backward pass accounts for majority of costs in model training.
- Backward pass has to compute input gradient & weight gradient both from output gradient.
- This incurs off-chip memory accesses, the most time consuming part in training.

Prior Work

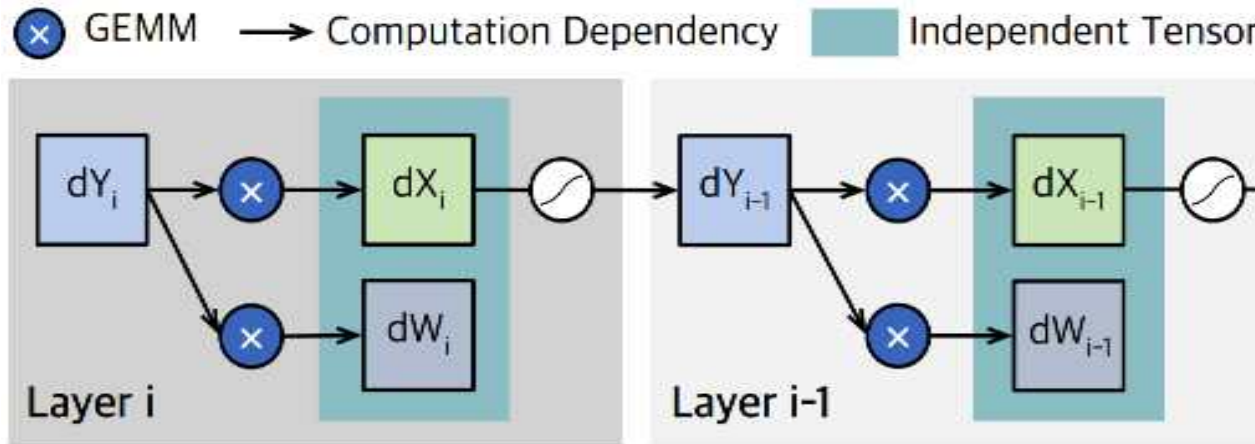
- Prior works focus on intra-operation data reuse.
- This cannot represent multiple independent operations, it can only represent single nested loop.
- Thus, they lose chance of performance improvement opportunity by data reuse.

Motivation : DNN Training Ratio



-> Backward pass occupies most of time on training process, indeed.

Motivation : Redundant Access

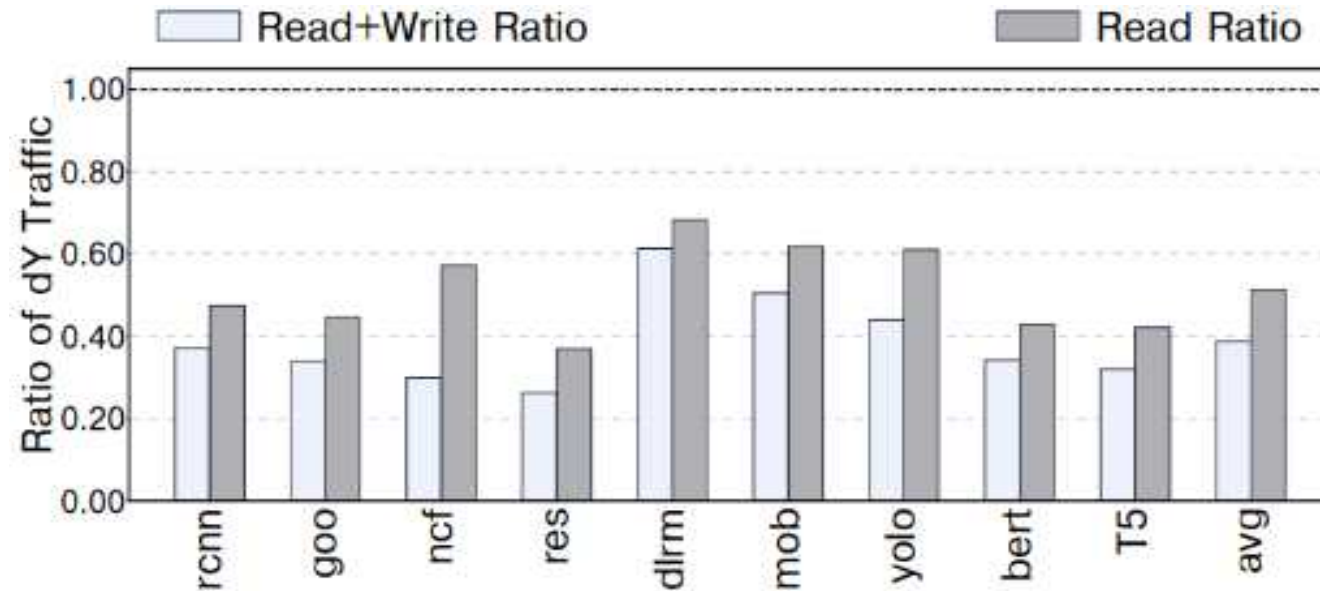


$$dX_i = \frac{\partial \mathcal{L}}{\partial X_i} = \frac{\partial \mathcal{L}}{\partial Y_i} \frac{\partial Y_i}{\partial X_i} = dY_i \times W_i^T$$

$$dW_i = \frac{\partial \mathcal{L}}{\partial W_i} = \frac{\partial Y_i}{\partial W_i} \frac{\partial \mathcal{L}}{\partial Y_i} = X_i^T \times dY_i$$

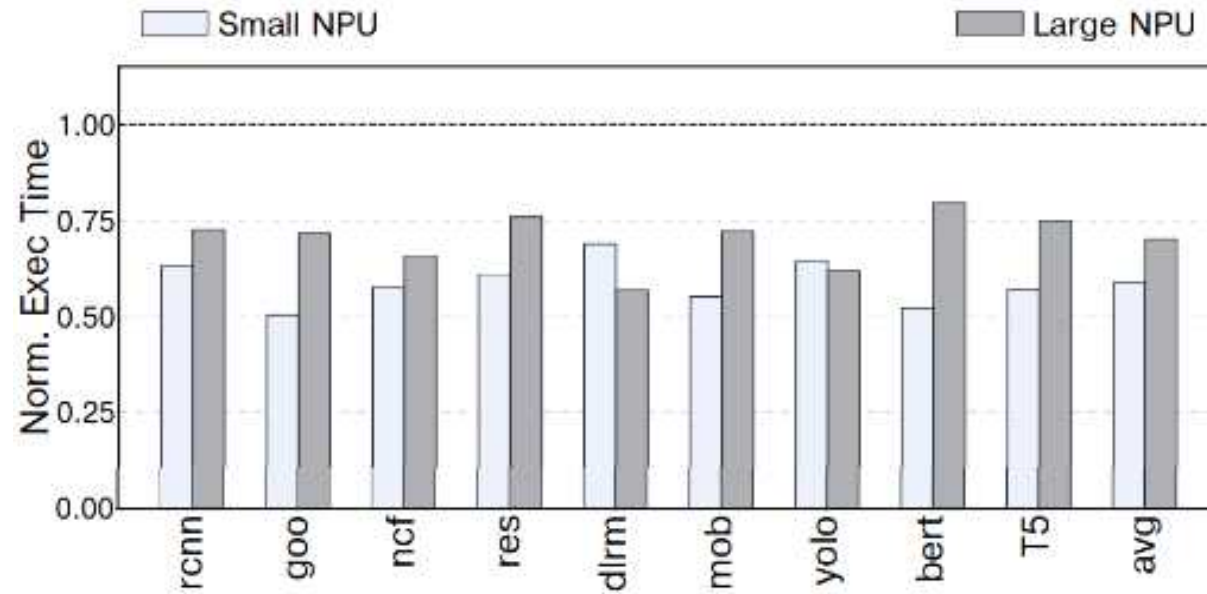
- In the figure above, dX_i and dW_i are computed by using the same dY_i as an operand.
- As these are independent, it is possible to reorder or even fuse two operations in an interleaved manner.
- Additionally, redundant accesses to the output gradient (dY_i) can be potentially reduced by such fusion.

Motivation : Data Traffic of dYi



- The ratio of dYi traffic compared to all read and write data, which is 39.0% of total data traffic on average.
- Furthermore, dYi occupies 51.4% of read data on average.
- This shows opportunities to improve performance by reducing the dYi traffic.

Motivation : Potential Performance

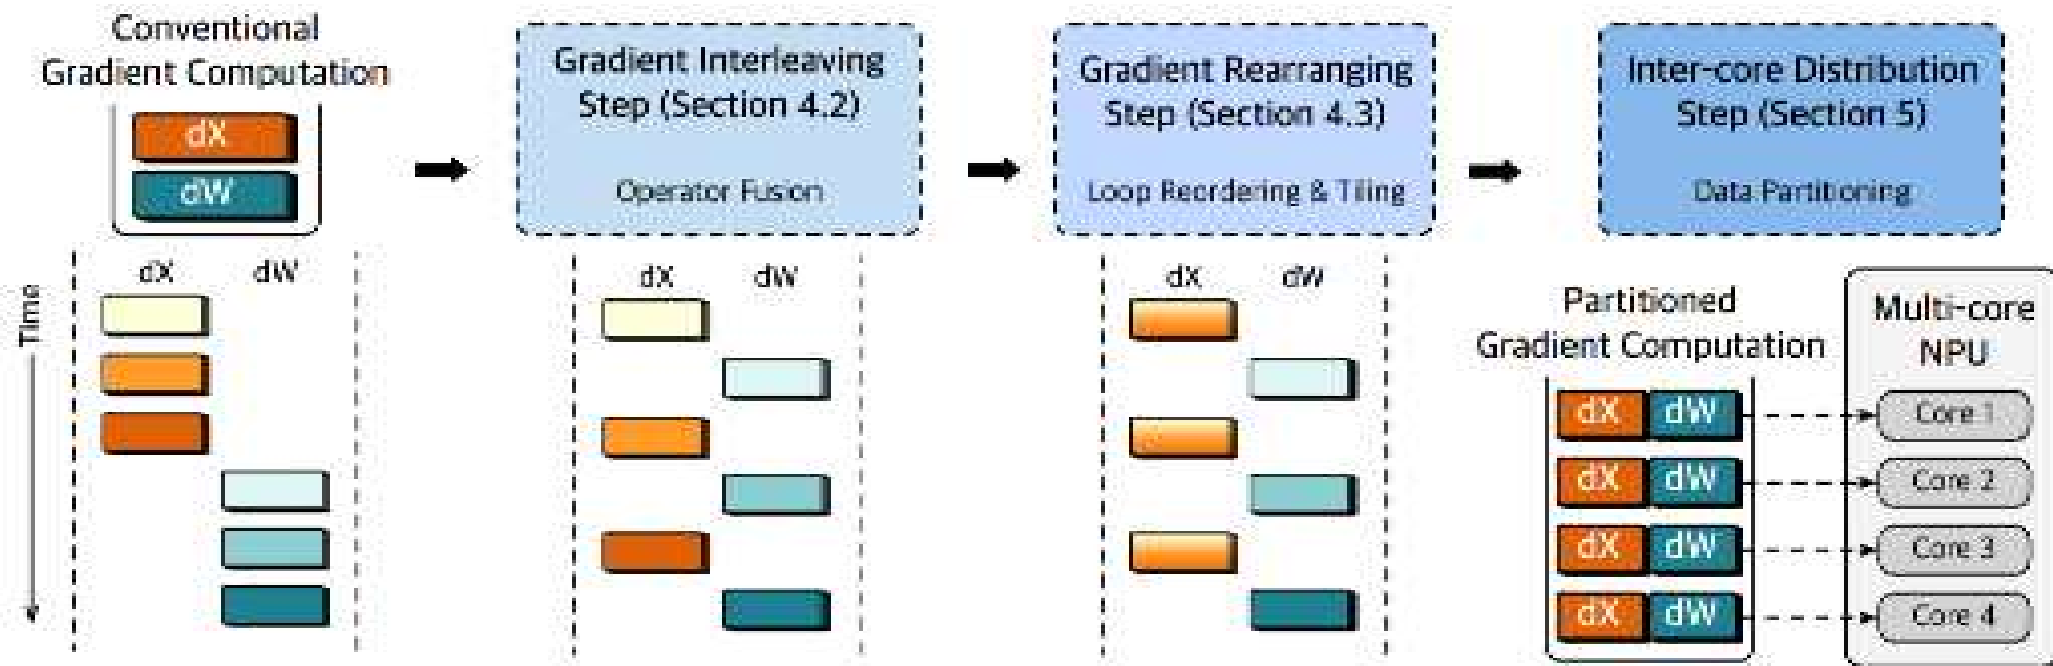


- How much is the performance potential of eliminating redundant reads for the output gradient?
- Researchers simulated with the elimination one of two accesses for dY.
- As a result, the speed up against the baseline is 1.43x in Large NPU and 1.70x in Small NPU.

Data Reuse Opportunity

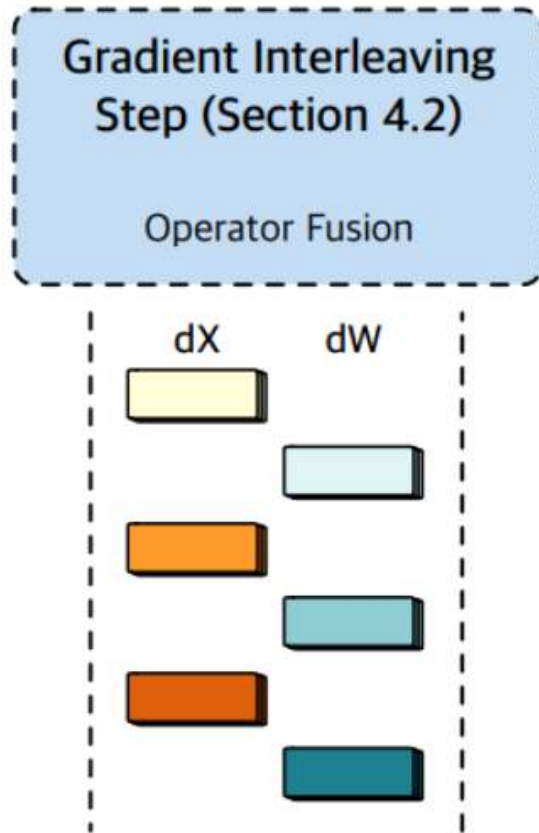
- Off-chip memory access is a problem → we need on-chip SPM, Scratchpad Memory.
- This paper introduces a novel code transformation technique aimed at removing unnecessary memory accesses for the output gradient.
- This is called interleaved gradient order, which enables data reuse on SPM.

Overview



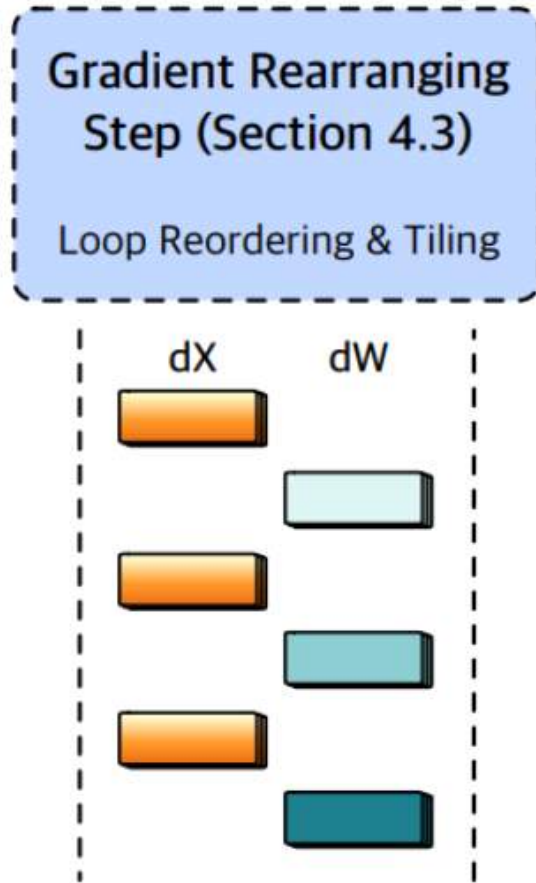
1. Gradient Interleaving Step
2. Gradient Rearranging Step
3. Inter-core Distribution Step

Overview : Gradient Interleaving Step



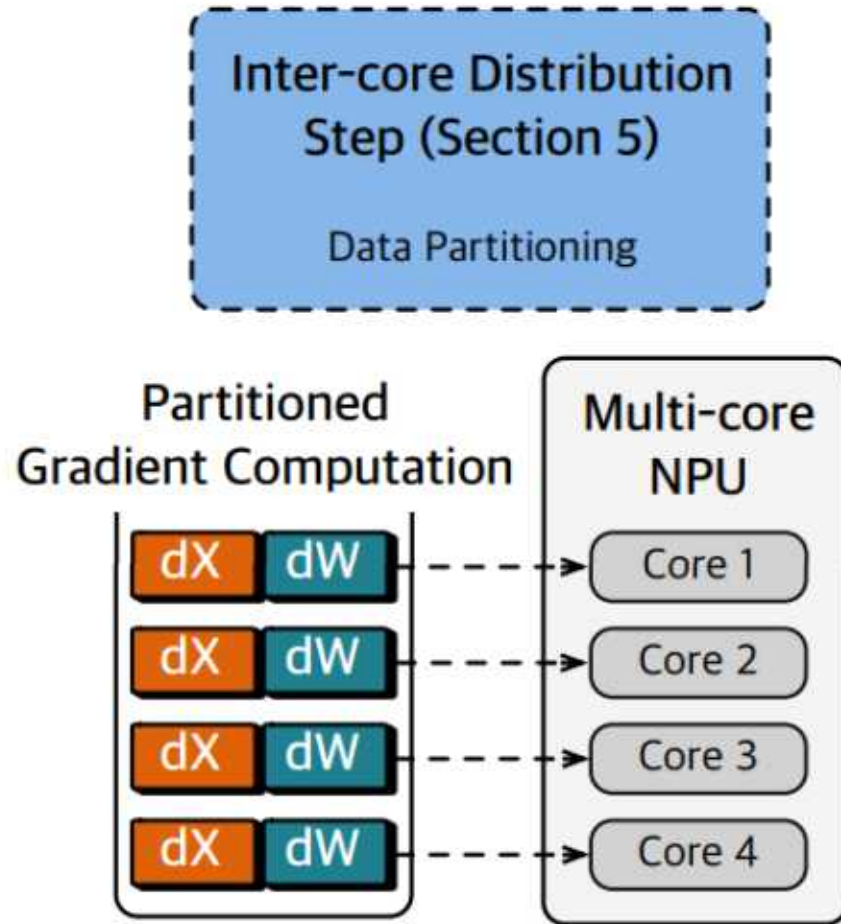
- This step combines two computations for dW and dX .
- Interleaves tiled operations for them.

Overview : Gradient Rearranging Step




- Effectiveness of tile reuses can vary depending on the operand dimensions.
- Thus, gradient rearranging step reorganizes interleaved dX and dW computations.

Overview : Inter-core Distribution Step




- For multi-core NPUs, operands are decomposed and assigned to different cores to maximize dY reuse.
- This step determines the optimal methods for operand decomposition.
- Finally, each segment is allocated to an NPU core.

Gradient Interleaving Step



Baseline approach


- dX and dW are calculated in a sequential manner.
- dY can be transferred to SPM twice due to prior loaded tile being evicted.



Interleaved computation


- dX and dW computations are interleaved (no dependency between them).
- Transformed code eliminates redundant accesses to dY tiles.

Gradient Interleaving Step



Baseline approach

- Yellow starred dY_i loaded twice.
- Because distance between dX_i and dW_i calculations exceeds the capacity of SPM.




Interleaved computation

- Yellow starred dY_i can be reused just one loading when interleaving is applied.
- Thus, interleaving significantly reduces data traffic and boosts the utilization of SPM.

Gradient Interleaving Step : Problems


- Interleaving method does not significantly improve performance in certain layers that contain **non-square tensors**.
- Why? The access patterns for dY_i differ between dX_i and dW_i computations.

Gradient Interleaving Step : Problems




- When computing dX_i , the access of dY_i follows a row-major access order.
- When computing dW_i , dY_i is accessed in a column-major order.


Three Groups of Memory Access Orders



- Traditional access order



- row-major access for both dX_i and dW_i



- column-major access for both dX_i and dW_i

Selection Algorithm

- + Data reuse of dY_i can be enhanced by appropriately altering the memory access order.
- Intermediate results can be generated in dW_{major} and dX_{major} due to the changed computation order.
- > Straightforward but fairly accurate algorithm is needed.

Selection Algorithm

```
1 GEMM in forward pass is  $X_i(M, K) \times W_i(K, N) \rightarrow Y_i(M, N)$   
2 if ALMOSTSQUARECOMPUTATION() then  
3   | Use Interleaving  
4 else if  $K > N$  and  $K > M$  then  
5   | Use Interleaving+dWmajor  
6 else  
7   | Use Interleaving+dXmajor
```

- Algorithm selects appropriate memory access order based on the shape of tensors.
- Nearly square tensors by *AlmostSquareComputation()* \rightarrow *Only Interleaving*.
- Non-square tensor
 - Column dim of $W_i >$ Row dim of $W_i \rightarrow$ *Interleaving + dWmajor*
 - Otherwise \rightarrow *Interleaving + dXmajor*

Gradient Rearranging Step

- Gradient rearranging step is combination of:
 1. Interleaved gradient computations
 2. Selection of optimal tile access order




Data Partitioning

- Single GEMM divided into smaller partitioned GEMMs.
- Performance depends on the dimensions of the GEMM.
- Thus, we need new data partitioning schemes to better suit interleaving and rearrangement steps.

Data Partitioning

Three Tensors:


- X
- W
- Y



- It is common for GEMM to be partitioned on a batch basis (dimension M).
- dY and X are split for computing dX and dW .

Data Partitioning : Problem

- Partitioned on a batch basis has a problem.
 - dW: dY being right-hand-side operand
 - dX: dY being left-hand-side operand
- > This prevents a single partition from performing all the necessary calculations.
-> Also it might have to accumulate intermediate results from multiple partitions.
-> If M is not significant in GEMM, this is not an efficient way.

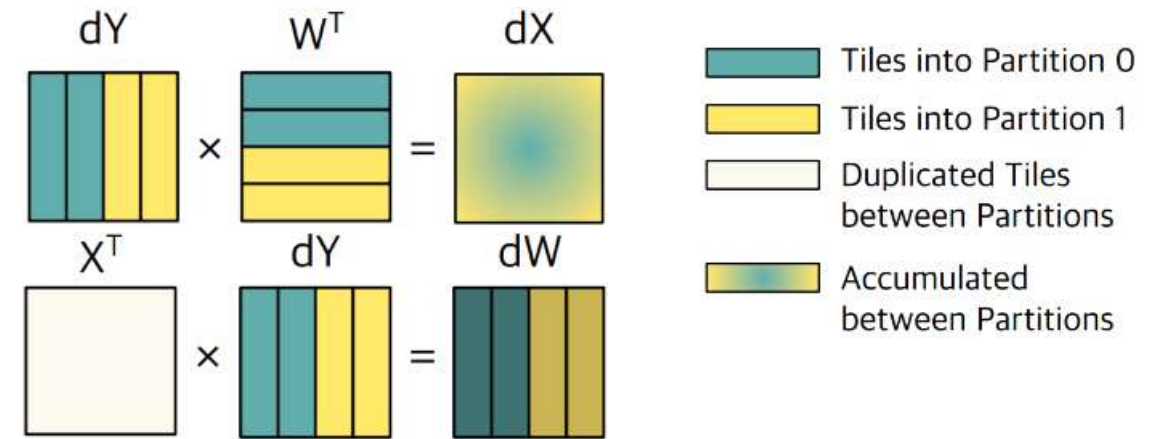


Data Partitioning in Another Dimensions


Partitioning GEMM in N or K dimension rather than M dimension could be more efficient in some layers.

Data Partitioning in Another Dimensions

- Partitioning GEMM in N
- All X is duplicated.
- Requires another accumulation



- Partitioning GEMM in K
- All dY is duplicated.
- Requires no accumulation




Data Partitioning : Selection

- Finally, overall performance depends on the dimensions of tensors.
- What if dimension M is significant rather than N and K in GEMM?

-> Systolic array can be fully utilized.

-> Because dimension M is wide enough. If it is not, other data partitioning scheme should be considered.



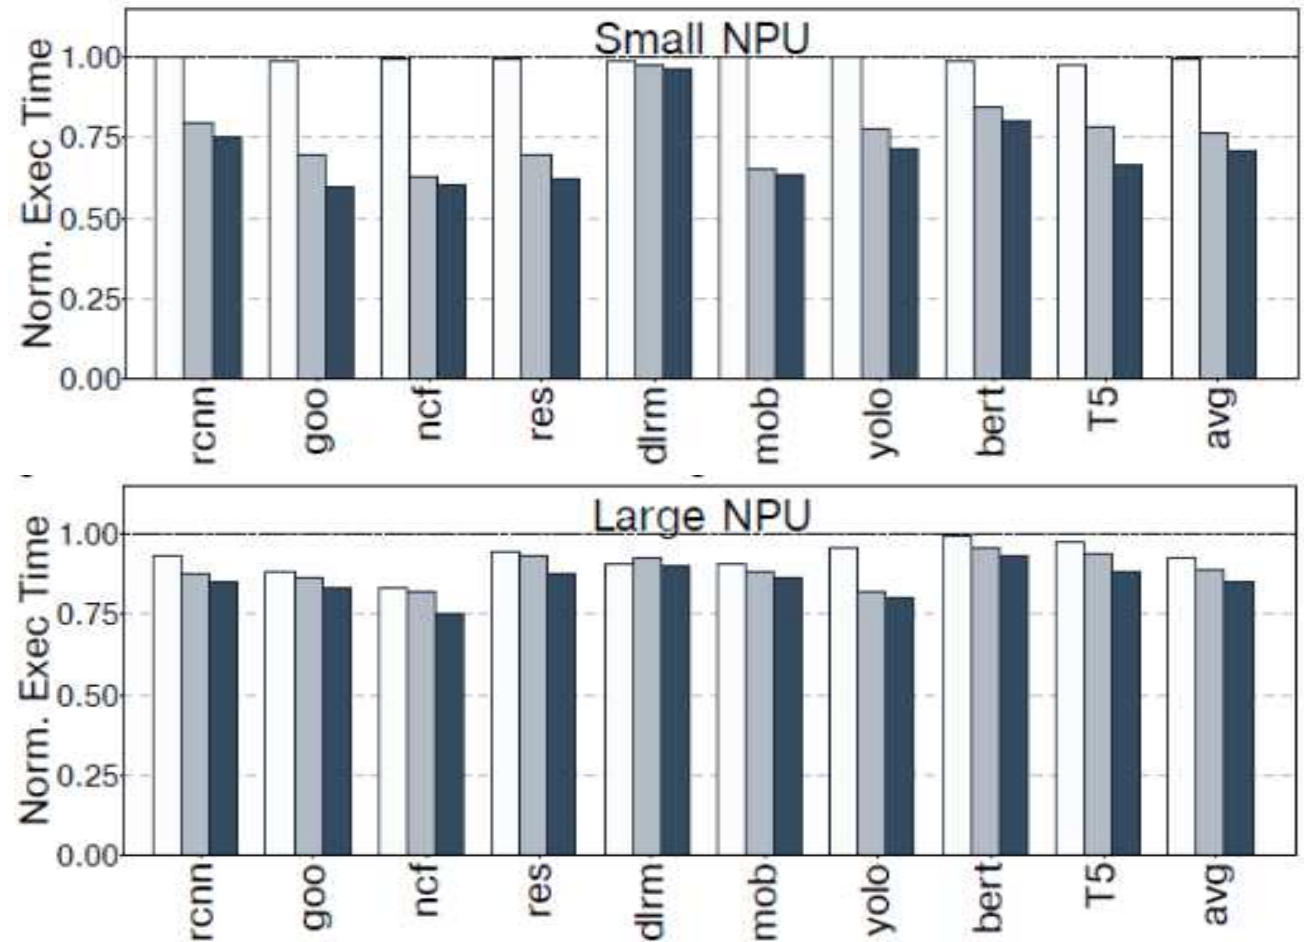
Data Partitioning : Selection

- We need to determine optimal data partitioning schemes for each layer.
- Thus, KNN algorithm is applied using the dimensions of dX , dW , dY as features.
- It has 91% accuracy and wrong predictions are negligible for performance.

Performance Evaluation

- Performance improvement on average
- Small NPU: 29.3%
- Large NPU : 14.5%

□ +Interleaving ▒ +Rearrangement ■ +DataPartitioning



Performance Evaluation

- Small NPU has more performance improvement since small NPU has smaller SPM.
- Thus, on-chip data reuse by interleaved gradient order becomes more significant.

Summary

- Interleaved Gradient Order maximizes data reuse within SPM during the backward pass of DNN training.
- Improves performance by inter-operation data sharing, not just intra-operation as prior works.
- This has three key steps.
 - 1) Interleaves gradient computation
 - 2) Select optimal tile access order
 - 3) Data partitioning by selection algorithm
- Performance improvement of 29.3% for small NPU and 14.5% for large NPU.